

THE DAILY & THE RECORD · 2026-08-20

******Filed verbatim from the secretary's first run (Donna, run-secretary-1, dispatched 2026-08-21 on the CEO's release "lets run her for yesterday"). CTO filing note at the end.******

TL;DR

...

The book fell 6.31% to \$1,884.79 and ended the day halted. Almost the whole loss was one bad fill, not the market: a test exit rule sold a real position at a fabricated \$100 for an asset that closed at \$415.04 — \$128.26 destroyed, and the drawdown reference it corrupted now caps the fund's risk capacity for twelve months. The bench confirmed twelve defects in the firm's own instruments the same day and the follow-on sells left the book clean: two positions, 26.6% invested, \$1,383.46 in cash. One thing needs you: choose how to repair the corrupted drawdown reference — repairing it moves the odds of another forced halt this year from 58.4% to 9.6%. Two approved buys totalling \$414.76 and the resumption of trading wait only on the clock.

...

THE DAILY · 2026-08-20

The book. NAV closed \$1,884.79 (seq 562, 23:46Z) against \$2,011.81 at the prior strike (seq 209, 2026-08-19T21:57Z) — **-\$127.02, -6.31%**; two positions (DBC \$252.68, TLT \$248.66), cash \$1,383.46 (73.40%), gross \$501.34 (26.60%), down from nine positions and 82.76% gross at the open; **trading HALTED** on daily loss since 15:52:31Z (seq 441, "Daily NAV loss 6.33% exceeds limit 4.00%").

What moved

- A **machinery-test** loss rule fired on a **fabricated \$100.00 GLD mark** and auto-policy v1 approved the sell (seq 255–258, 08:01Z) — realised destruction **\$128.26** against a \$170.71 basis, ~85% of the fund's 7.42% drawdown, and the event that fired the halt (docs/pm/PM_SLEEVE_V2_2026-08-21.md §3).
- Eight sells filled for **\$1,036.54** — one auto-approved, three by **rushi**, four by **rushi-via-cto** on the CEO's quoted remote instruction (seq 258, 286, 289, 292, 377, 380, 383, 386); three legacy strategies were then archived flat (seq 387–389).
- The bench delivered **16 dispatches** and confirmed **twelve defects** in the fund's own instruments, including one KILL (run-adversary-v5r3, gate v5 round 3) and a register limit whose discrimination runs backwards (run-validator-r6d2).

- The COO ran twice, compressing **20 open items** → **6 batches** and then **73 items** → **10 truly open** → **7 batches** (run-coo-1, run-coo-2); the CEO decided **103 recommendations** — 101 accepted, 2 rejected.
- Two seats were created — **coo (Vishesh)** at 15:50Z (commit f957050) and **secretary (Donna)** at 20:25Z (commit bf64096) — and approval identity migrated [rushy → neesh](#) in guard v1.1 (commit f9171e7).

Awaiting you

1. **PM R1 — the corrupted drawdown reference** (seq 502): repair by correction event, versioned dated exception, or accept and run reduced for twelve months. Repairing moves zero-drift P(halt in 252d at 61% gross) **58.4% → 9.6%**. It gates the sleeve v2 design.
2. **Halt resume** — you took the no-decision (seq 403, 510); \$1,383.46 immobilised until the loss reference rolls at 2026-08-21T00:00Z.
3. **The two staged buys** — SPY \$263.94 and DBA \$150.82 (seq 546, 547), queued behind that clock.
4. **Gate v5 round 4** — filed 19:47Z (commit abe36d2), awaiting the blind adversary; no attack resolved against it on 08-20.
5. **COO Batch E** — the only batch you rejected (seq 508); the record does not state what replaced it.

You approved 103 recommendations in six sittings today and turned down two — the last thirty-four took six minutes fifty-six seconds. That is either total command of your own firm or a very fast index finger, and the log politely declines to say which.

THE RECORD · 2026-08-20

NAV & THE BOOK

	OPEN (SEQ 228, 00:26Z)	CLOSE (SEQ 562, 23:46Z)	Δ
NAV	\$2,011.81	\$1,884.79	−\$127.02 (−6.31%)
Positions	9	2	−7
Cash	\$346.92 (17.24%)	\$1,383.46 (73.40%)	+\$1,036.54
Gross exposure	\$1,664.89 (82.76%)	\$501.34 (26.60%)	−\$1,163.55

Opening book: SOFI 9.1882 @ 18.42, INTC 1.6088 @ 92.80, MSFT 0.3401 @ 484.31, SPY 0.2178 @ 769.06, GLD 0.4245 @ 413.84, NVDA 0.7499 @ 217.56, XLE 2.7499 @ 63.58, DBC 8.1222 @ 30.76, TLT 3.0199 @ 83.02 (seq 228). Closing book: DBC 8.122157 @ 31.11 (\$252.68) and TLT 3.019871 @ 82.34 (\$248.66), both inside [sleeve_beta_500](#) and both covered by pre-committed loss and time rules (seq 355, 356).

Sixteen NAV strikes were taken during the day. The prior day's last strike also read \$2,011.81 (seq 209), so the −6.31% is measured against an unmoved reference. Drawdown at the close read **7.74% max, 7.44% current** against a trailing-365-day peak of \$2,036.35 and a 10% limit ([GET /fund/risk/monitor](#)); the peak has never been rebased, and §VI D1 records why that matters. The PM read NAV \$1,885.25 at 19:15–19:16Z (docs/pm/PM_SLEEVE_V2_2026-08-21.md §1), a different strike, not a disagreement.

Trading halted eight times and resumed seven; every resume was auto-reverted within roughly one monitor tick because the daily-loss condition had not cleared. The book therefore ended the day halted.

#	HALT (Z)	LOSS CITED	RESUME (Z)	BY
1	08:16:08 (seq 265)	6.62%	08:33:34 (seq 281)	rushi
2	08:33:43 (seq 282)	6.62%	09:05:17 (seq 297)	rushi
3	09:05:32 (seq 298)	6.62%	14:24:51 (seq 396)	neelesh
4	14:26:02 (seq 400)	6.35%	14:35:39 (seq 429)	neelesh
5	14:35:48 (seq 430)	6.36%	15:14:35 (seq 433)	neelesh
6	15:14:42 (seq 434)	6.31%	15:21:51 (seq 435)	neelesh
7	15:22:01 (seq 436)	6.30%	15:52:30 (seq 440)	neelesh
8	15:52:31 (seq 441)	6.33%	—	still halted at close

Leg 3 of the team metric — capital deployed under mandate. The book closed at 26.60% gross against a normal gross of 61% adopted the same evening (seq 499) and a throttled target of 48.6% (run-pm-sleeve-v2). The PM records the gap as a leg-3 defect: **\$1,289.20 idles above the \$94.26 cash floor** (docs/pm/PM_SLEEVE_V2_2026-08-21.md §2). The two staged buys, \$263.94 + \$150.82 = \$414.76, would carry gross to 48.60% (seq 546, 547).

|| TRADING & EXECUTION

Eight fills, all SELL, all paper venue, \$1,036.54 total notional. The approver column is the day's governance record: one auto-policy approval that should not have happened, three under the operator's own name, four through the via-CTO channel carrying the CEO's quoted instruction.

SYMBOL	SIDE	QTY	PRICE	VENUE	APPROVER	SEQ
GLD	sell	0.424471	\$100.00	paper	auto-policy-v1	256 / 258
SPY	sell	0.052217	\$769.0599975585938	paper	rushi	284 / 286
INTC	sell	1.608762	\$92.80000305175781	paper	rushi	287 / 289
MSFT	sell	0.340051	\$484.30999755859375	paper	rushi	290 / 292
NVDA	sell	0.749886	\$219.60000610351562	paper	rushi-via-cto [CEO (Neelesh), remote]	375 / 377
XLE	sell	2.749912	\$64.63500213623047	paper	rushi-via-cto [CEO (Neelesh), remote]	378 / 380
SPY	sell	0.16554	\$767.1900024414062	paper	rushi-via-cto [CEO (Neelesh), remote]	381 / 383
SOFI	sell	9.18819	\$18.559999465942383	paper	rushi-via-cto [CEO (Neelesh), remote]	384 / 386

Seventeen orders were proposed and nine declined. The declines are the day's second governance story: one by the CTO's stale guard — "the GLD position this ticket closes was already sold by the phantom-price exit fire at 08:01Z; approving now would SHORT 0.424471 GLD" (seq 283); four re-staged by the CTO after ageing past the 120-minute freshness window (seq 323–326); four expired automatically by the worker at 12:21:55Z on the same rule (seq 364–367).

Three approvals were **refused** by the approval-channel guard: the pm seat attempting an order approval (seq 373, "approver 'pm' is not on the approval allowlist"), a missing confirm echo (seq 374, "approving this order requires confirm='20ab9a86'",) and rushi against the migrated desk-request allowlist (seq 395). Note for the record without resolving it: at 13:42Z the order allowlist read ['rushi', 'rushi-via-cto'] (seq 373) while at 14:07Z the desk-request allowlist read ['ceo', 'neelesh', 'neelesh-via-cto'] (seq 395) — the two surfaces migrated at different moments; commit f9171e7 (19:37Z) records guard v1.1 as the migration.

One order was administratively terminated: INTC 3.0 sell, "the execution step died between OrderApproved (seq 74, 2026-08-14T13:20:04Z) and venue submission — no OrderSubmitted event and no venue_ref exist" (seq 560, written 2026-08-20T21:39:27Z).



Sixteen dispatches resolved, every one on Opus. Verdicts quoted as recorded in `fund_agent_runs`.

RUN	SEAT	RESOLVED (Z)	VERDICT (QUOTED)	ARTIFACT
run-pm-staging	pm	03:30	"8 SELL TICKETS + 3 STATE TICKETS + 7 RECOMMENDATIONS"	docs/pm/PM_STAGING_2026-08-20.md
run-builder-1	builder	04:07	"MERGED — tsc clean, 17/17 tests, 8 seat routes 200, nav 5 tabs"	KryptonPay 25f8e61..d8c2820
run-riskofficer-1	riskofficer	09:04	"POLICY CORRECT, WORLD FALSE; FIX INCOMPLETE (F1); PREMISE FALSE (F2); MARKER FORGEABLE (F3)"	docs/AUDIT_AUTOPOLICY_V1_FIRST_FIRE_2026-08-20.md
run-builder-2	builder	10:03	"SHIPPED (2 commits, tsc clean, 37/37, build green) + 11 KryptonPay findings + 10 ClarkHarness findings + 7-screen design audit; 1 CRIT live defect confirmed against the spine"	KryptonPay 406434e..60ba593
run-adversary-v5r3	adversary	10:36	"KILL"	docs/reviews/ADVERSARY_GATE_V5_R3_2026-08-20.md
run-coo-1	coo	10:52	"20 OPEN ITEMS → 6 BATCHES + 1 NON-DECISION; 2 OBJECTIONS (1 constitution violation, 1 funnel ordering); 2 items re-scoped"	docs/coo/TRIAGE_2026-08-20.md
run-mechanism-cycle1	mechanism	11:23	"0 PROPOSALS TO THE CONTAINER: entry 5 RETIRED (arithmetic), entry 6 DECLINED (mechanism falsified on our feed), entry 11 SPEC-FILED/DEFERRED (2 prerequisites); 3 instrument defects found"	docs/research/MECHANISM_CYCLE1_2026-08-20.md
run-builder-dispatch3	builder	13:52	"DELIVERED: 4 commits, 26 files +2060/-257 ... 80/80 tests ... SIX absence defects fixed, three of them in code written this dispatch"	KryptonPay a528396
run-validator-r6d2	validator	14:24	"THREE FINDINGS... max_risk_concentration_pct=0.50 has INVERTED discrimination... NEITHER limit gates anything (display-only)"	docs/AUDIT_R6_D2_ATTRIBUTION_2026-08-20.md
run-cdo-trial-1	cdo-trial	14:43	"DELIVERED — 10 ranked defects (2 money-visible...), design language made explicit and versioned, floor spec with 7 falsifiable acceptance criteria"	KryptonPay/docs/design/CDO_AUDIT_2026-08-20.md
run-validator-wfgeom	validator	15:00	"INVERSION CONFIRMED with two exact closed forms...; NO issued verdict affected (prospective only)... window_for's fold count is INVARIANT to available history"	docs/AUDIT_WF_GEOMETRY_2026-08-20.md

RUN	SEAT	RESOLVED (Z)	VERDICT (QUOTED)	ARTIFACT
run-builder-dispatch4	builder	18:54	"DELIVERED A, B, C, D... ClarkHarness 961 → 1034 tests, KryptonPay 80 → 94... Every part landed EXCEPT the root cause of F4's 14m41s halt latency"	ClarkHarness builder-d4 merge (base bd4b30c)
run-coo-2	coo	19:08	"73 ITEMS → 10 TRULY OPEN → 7 BATCHES... trigger... miscounting by 3.65x... 1 SELF-REVERSAL... 3 controls found without a working surface"	docs/coo/TRIAGE2_2026-08-20.md
run-pm-sleeve-v2	pm	19:40	"DESIGN DELIVERED — four legs, normal gross 61%, throttled to 48.6% today... GATED ON R1"	docs/pm/PM_SLEEVE_V2_2026-08-21.md
run-quant-entry11	quant	20:13	"THREE FAILS, BAND DOES NOT FLIP ON COST... the one genuine kill is 2-of-4 fold persistence... Hybrid draft trial NOT confirmed"	docs/research/QUANT_ENTRY11_2026-08-21.md
run-builder-dispatch5	builder	21:41	"DELIVERED A, B, F, E, H — 1100 ClarkHarness + 127 KryptonPay passing... STOPPED at the sanctioned boundary: C, D, G NOT built"	ClarkHarness ea53865..b490ca0 + KryptonPay 2bb3095..24a12dd

Leg 2 of the team metric — candidates reaching the belt. One: entry 11, run on the belt at slippage {1, 3, 5} bps and failing gate v4.1 at all three (run-quant-entry11). The mechanism's funnel cycle produced **zero** new proposals — two honest refusals and one deferred spec (run-mechanism-cycle1, commit 8c313d2, filed as *"o proposals, and that is the machine working"*). Against the funnel doc's 3–5/week target, the day contributed one belt run and no new candidates.

IV DECISIONS & GOVERNANCE

188 [DeskRecommendationDecided](#) events were written: 103 by the CEO (101 accepted, 2 rejected) and 85 by the CTO recording execution (73 done, 12 staged). The versioned and structural decisions:

DECISION	DECIDED BY	WRITTEN REASON, AS RECORDED
Auto-approval envelope v2	CEO → CTO (commit 441d04c)	"the envelope now verifies the number and the rule, not just the order" — riskofficer R1/R3/R4 implemented (seq 331–333)
Auto-approval envelope v3	CEO (seq 405), commit ac53f31	"cut envelope v3 = v2 + riskofficer R5... State plainly R5's blast radius today is \$0"
Approval-channel guard v1 , then v1.1	CTO (68ba29d), CEO instruction (f9171e7)	"approval identity migrates rushi → neesh (CEO instruction)"
Equity seed prices deleted	CEO (seq 335), commit e8223f2	"Ratified; seeds deleted" — "a stale chosen number is a fabricated mark"
Paper venue may never fabricate a price	CTO, commit fa6b877	"Incident fix: the paper venue never fabricates a price again"
machinery-test GLD loss rule overridden	CEO via COO batch B, staged by CTO (seq 354)	"test artifact, not portfolio policy — created 2026-08-17 to prove the exit wiring, fired 2026-08-20 on a fabricated mark"
TLT and DBC thesis exit rules set	CEO-decision-staged-by-cto (seq 355, 356)	"Held per the RECORDED sleeve fact approved by the CEO 2026-08-20... declared beta, no edge claimed"
Three legacy strategies retired	CEO (PM recs 2/3/4), CTO executed (seq 387–389)	"CEO-accepted retirement, book flat as of 2026-08-20T13:48Z" — Trend · Sector & Commodity; Momentum · Large Cap Tech; Mean Reversion · Cyclical
Concentration limit re-specified	CEO (seq 410), commit bd4b30c	"max_risk_concentration_pct = 0.50: RETIRE or re-specify" → "component vol replaces risk share"
Sleeve v2 normal gross = 61% of NAV	CEO (seq 499)	"derived by bootstrap, argued in the memo"
Target gross = normal × throttle	CEO (seq 524)	"/risk/throttle becomes the sizing input"
BUY SPY \$263.94 into new sleeve_premia_equity	CEO (seq 525), staged (seq 546)	"stop 7.3% and time +90d committed BEFORE the order"
BUY DBA \$150.82 into new sleeve_premia_carry	CEO (seq 521), staged (seq 547)	"stop 6.1% and time +90d committed first; ag-contango caveat stated"
Halt resume — no decision	CEO (seq 403, 510)	"no halt decision tonight; resume after 2026-08-21T00:00Z when the loss reference rolls and the alarm clears by arithmetic. OBJECT to using the rebase route to skip the wait."
Gate v5: no total-book PSR floor on benchmarked claims	CEO (seq 551)	"the criterion rejects on beta while claiming to test luck"
Gate v5: breakeven measured on the excess	CEO (seq 555)	"total return dies at 28.1bps while the alpha dies at 5.92bps, and only the second decides an alpha claim"
Quant hybrid split reverted for whole algorithms	CEO (seq 557, *rejected*)	"the split as designed is reverted by this trial's measurement"
COO Batch E rejected	CEO (seq 508)	the batch as framed — decide the colour item once, defer the 3D floor, adopt or revise the reversibility table — was declined; the record does not state the substitute
First seat-filed dispatch ask blessed	CEO (seq 408, 14:29Z)	the mechanism's request for a validator fold-geometry audit (seq 360), fired by the CTO at 14:44Z (seq 431)
Seat coo (Vishesh) created	CEO, commits f957050 / 4c7af6d / 0484115	"the CEO desk triaged into batch decisions"; the seat "endorses, never decides"
Seat secretary (Donna) created	CEO, commits bf64096 / 9bcd89b /	"the EoD record (CEO decision)"; the observer's note added at 23:42Z

DECISION	DECIDED BY	WRITTEN REASON, AS RECORDED
	c5d0194 / a00249c	
Batch-by-seat default + COO trigger at >20 desk items	CEO instruction, commit 52626cd	dispatch rules amended the same day the desk overflowed
Seat-filed dispatch requests are asks, never triggers	CEO, commit 56bdc69	"that design looks good" — the org chart gains edges, the ignition keys stay human

V INSTRUMENTS & INFRASTRUCTURE

Ninety-nine commits landed on 2026-08-20 UTC — **34 ClarkHarness, 35 KryptonPay, 30 firm** (method: committer date in UTC across all three repos).

The gate moved twice: **v4.1** shipped at 08:42Z (commit cc11c6f, "the holdout leg gets the discipline the folds always had", 931 tests passing per seq 231), then **round 3 was killed** by the blind adversary and the kill was filed rather than argued with (commit ffda1a8, "the third kill in the chain"), and **round 4** was rebuilt and filed at 19:47Z awaiting review (commits 34ebcf8, abe36d2). The auto-approval envelope went v1 → v2 → v3 in a single day (441d04c, ac53f31), each version a response to a finding from the seat that supervises it. The flight recorder went in at 04:26Z and gained reasoning and trace threading half an hour later (5a24ed3, b1f415e), which is why this record can be written at all.

Test suites, as recorded by the builder at each merge: ClarkHarness **931** (gate v4.1) → **961** → **1,034** (dispatch 4) → **1,100** (dispatch 5). KryptonPay **17/17** (dispatch 1) → **37/37** (dispatch 2) → **80/80**, 37 inherited plus 43 new (dispatch 3) → **80** → **94** (dispatch 4) → **127** (dispatch 5). Dispatch 5 also verified the office floor **live** — "14 real pulses of 247 events, halted-room state correct" — and stopped at the sanctioned boundary with parts C, D and G unbuilt, on the builder's own reasoning that "a two-dispatch feature built well costs less than one built twice."

Also shipped: the archive PDF pipeline and letterhead this memo will be rendered through (commits 73dd214, 3c66aa5, bfoea6b), the spine's attribution guard, two-sided cap, correction event and halt classes (9eb4d11), venue-on-fills, informative TCA cuts, **money_at_stake** and the desk-load counter (838fc9a), the generation machine — premia menu, weekly funnel rhythm, revival register (64ed416) — and the cost model (fec38d4).

Seventeen dated documents were filed: three audits (AUTOPOLICY_V1_FIRST_FIRE, R6_D2_ATTRIBUTION, WF_GEOMETRY), the GLD incident, gate v5 rounds 3 and 4, the min-train-return review, the cost model, two COO triages, the PM staging memo, four research documents, and two adversary reviews.

VI THE DEFECTS LEDGER

Twelve confirmed defects in the fund's own beliefs, instruments and code. Ten were found by seats; one by the CTO's stale guard in the moment; one by the deterministic monitor.

#	DEFECT (CONFIRMED)	FOUND BY	MONEY IT COULD HAVE TOUCHED
D1	The paper venue fabricated a \$100.00 mark for GLD against a \$415.04 close; a test loss rule fired on it and auto-policy approved the sell	monitor → riskofficer (run-riskofficer-1), PM (run-pm-sleeve-v2)	\$128.26 realised. ~85% of the 7.42% drawdown; absent it, ~1.12%. The corrupted \$2,036.35 peak is trailing-365-day, so it caps risk capacity for twelve months; repairing it moves P(halt in 252d at 61% gross) 58.4% → 9.6% (seq 502)
D2	Exit enforcement sizes by symbol, ignoring strategy_id (<code>exitrule.py:269-270</code>) — machinery-test's rule sold Trend's GLD	pm, CTO-verified line-exact	The mechanism of D1. Design rule adopted: no symbol in two strategy_ids
D3	Auto-policy v1: fix incomplete (F1), premise false (F2), marker forgeable (F3) — "POLICY CORRECT, WORLD FALSE"	riskofficer (run-riskofficer-1)	The whole auto-approval surface. Closed by v2 (441d04c) and v3 (ac53f31); both audit forgeries now fail closed (seq 332)
D4	One mistagged fill produced a phantom ± \$174.47 pair: GLD bought under Trend (seq 76 @ \$402.18), sold under machinery-test (seq 258)	validator (run-validator-r6d2 ITEM 3)	A rebalance preview returned a \$376.84 GLD BUY on a book holding zero GLD , bounded at \$753.68 (40% of NAV) . Corrected by <code>StrategyAttributionCorrected</code> (seq 444)
D5	<code>max_risk_concentration_pct = 0.50</code> has inverted discrimination — a 100%-single-name book reads 100.00%, *lower* than the hedged book's 102.49%; and neither structural limit gates anything (display-only)	validator (run-validator-r6d2)	A concentration control the register believed it had. Re-specified as component vol (commit bd4b30c)
D6	<code>/fund/tca</code> emitted " reliable: true, -12.59bps, cheaper than modelled " — approval-latency drift on a venue where execution slippage is identically zero	validator (run-validator-r6d2)	Would have certified the 5bps cost constant as validated. Verdict leg fixed at resolve (commit 6caec4a)
D7	Walk-forward fold count is invariant to available history — GATE_V5_ROUND4 §3's history table measured a fold generator the belt does not implement; the 10y backfill buys zero folds	validator (run-validator-wfgeom)	A gate round's headline argument. Corrected by new section §8 rather than an edit (commit bf0ea6b)
D8	Fold count and regime coverage invert for short holds: <code>span_oos = K.floor(4h.365/252)</code> , a 24.2× hold-21 vs hold-1 span ratio at every K	mechanism D1 → validator confirmed	No issued verdict affected (prospective only); coverage qualifier accepted (seq 511)
D9	<code>max_component_vol_pct</code> is scale-invariant — moving gross 26.6% → 61% does not move it one basis point; <code>/risk/advanced</code> prints ES \$32.48 where the NAV-dollar figure is \$8.65 (3.76× overstated)	pm (run-pm-sleeve-v2 E3), CTO-verified	A composition control being read as a size control. Advisory-only; no money moved on it
D10	<code>desk_load</code> counted accepted+staged items — the COO trigger miscounted by 3.65× (73 against a live reading of 16)	coo (run-coo-2), on its own dispatch trigger	\$0 direct; a governance trigger firing on a false number. Fixed by the CTO the same hour
D11	Gate v5 round 3: KILL — the third kill in the chain	adversary, blind (run-adversary-v5r3)	The whole gate rewrite. Filed as a verdict, not argued (commit ffda1a8)
D12	An INTC sell sat 152 hours between <code>OrderApproved</code> (seq 74) and venue submission with no <code>OrderSubmitted</code> and no <code>venue_ref</code>	CTO reconstruction (seq 560)	"A late execution would have created an unintended short" — the position was already flat

Three further honest negatives belong here because they cost real work and returned no feature. The COO **reversed itself in writing**: sleeve falsification #2 is **not** discharged, because the four Batch-A fills were paper-venue and therefore zero-information by construction (run-coo-2). The mechanism returned **zero proposals** on three menu entries and said why for each — entry 5 retired on arithmetic (required alpha 16.9%/yr at top-5 venue costs, seq 428), entry 6 declined because the mechanism was falsified on our own feed, entry 11 spec-filed against two prerequisites. And the quant reported its own **hybrid trial failed**: the local draft was discarded on four harness-knowledge defects and reviewing it cost more than writing the file — which the CEO then used to revert the placement rule (seq 557). Each of these is leg 1 of the metric working exactly as designed, in service of leg 3.

VII THE FLOOR

Sixteen dispatches resolved on 2026-08-20 UTC, every one on Opus, totalling **3,216,188 tokens and 1,497 tool uses** — summed over exactly those sixteen runs, with no null token fields.

SEAT	RUNS	TOKENS	SHARE
builder	5	1,846,291	57.4%
pm	2	265,323	8.2%
validator	2	259,616	8.1%
coo	2	211,633	6.6%
cdo-trial	1	181,609	5.6%
quant	1	154,591	4.8%
mechanism	1	112,207	3.5%
adversary	1	93,131	2.9%
riskofficer	1	91,787	2.9%
Total	16	3,216,188	100%

No dollar cost is recorded per run; the log stores tokens and tool uses only. Applying the cost model filed the same day — ~\$0.70 per 100k tokens at Opus rates on a 90/10 input/output assumption, with an all-input floor of \$0.50 and an all-output ceiling of \$2.50 (docs/COST_MODEL_2026-08-20.md) — the day's bench work models to **≈\$22.51, bounded \$16.08–\$80.40**. That is a modelled figure, not a measurement: the input/output split is not captured, and the document says so itself.

New seats. `coo`, carrying the name **Vishesh** and a place in the executive row (commits f957050, 4c7af6d, 0484115); `secretary`, carrying the name **Donna**, created at 20:25Z with the persona and format contract following at 20:34Z and 20:43Z and the observer's note at 23:42Z (commits bf64096, 9bcd89b, c5d0194, a00249c). A **CDO was auditioned**, not seated — a trial brief, one artifact, and a seat memory recording the audition (commits 0a22202, e110b4c, 3e62e1d).

Workflow lessons that graduated. Batch-by-seat as the default and the >20-item COO trigger were written into the constitution (52626cd). Seat-filed dispatch requests became a durable, human-fired ask (56bdc69), and the first one was blessed the same day (seq 408). A shared API card was created so seats stop re-learning the spine by trial (f1117a2), and corrected twice within hours as measurements landed (ab8cd5e, and the cost gotchas in dofd1f). The validator's mandate gained a standing GAPS section aimed at generation (c93f2d2). The riskofficer's scope widened to the human approval channel (a1565fb) — the seat that supervises the machine now also supervises the people.

VIII CARRIED FORWARD

— **The halt.** Trading closed halted (seq 441). The accepted disposition is to resume after 2026-08-21T00:00Z when the loss reference rolls and the alarm clears by arithmetic, with a written objection to

using the rebase route to skip the wait (seq 510). \$1,383.46 immobilised.

- **PM R1 gates the sleeve.** The corrupted drawdown reference is the first thing tomorrow's session looks at; nothing in sleeve v2 sizes until it is decided (seq 502).
- **Two staged buys** wait on that clock: SPY \$263.94 into `sleeve_premia_equity`, DBA \$150.82 into `sleeve_premia_carry`, both with stops and 90-day time exits committed before the order (seq 546, 547).
- **Gate v5 round 4** stood filed and awaiting blind adversary review at the close (commit abe36d2, 19:47Z); no adversary run resolved against it on 2026-08-20.
- **Builder parts C (merge gate), D (factor pack) and G (Lab analytics)** were deliberately not built (run-builder-dispatch5), and **F4's 14m41s halt-latency root cause remains OPEN** by the builder's refusal to record it as fixed (run-builder-dispatch4, ratified at seq 515: "Do not record F4 as resolved").
- **Two PM-filed dispatch asks** sit in the queue awaiting a human trigger: mechanism on premia-menu entry 7, and mechanism/quant on trend-as-premia entry 4 (seq 548, 549, 20:08Z).
- **A dating note.** The north-star three-leg metric is recorded in the firm repo at commit 5521d48, committer-dated **2026-08-21T00:45:11Z** — after this record closes. It is noted here as carried-forward state, not as a 2026-08-20 decision, because the record dates it the following day.

IX THE OBSERVER'S NOTE

One. The CEO decided 103 recommendations in six sittings. The two heaviest were **34 items in 6 minutes 56 seconds** (seq 499–532, 19:54:56→20:01:52Z) and **28 items in 6 minutes 55 seconds** (seq 399–428, 14:25:52→14:32:47Z). The >20-item COO trigger exists for precisely this load, and it worked — run-coo-2 compressed 73 desk items to 10 truly open and 54 to bookkeeping. But it fired *after* the desk had filled, twice in one day, and its own counter was wrong by 3.65× when it fired. The instrument that protects the CEO's attention was, on the day it was needed most, both late and miscalibrated. That deserves someone's attention before the next twenty-item evening.

Two. One seat consumed the floor. `builder` ran **5 of 16 dispatches and 57.4% of the day's 3,216,188 tokens** — a mean of 369,258 per run against 124,536 for every other seat combined. The cost model filed that same morning put the central figure at "∼100k tokens per dispatch, range 67–122k, remarkably stable across seat types" (docs/COST_MODEL_2026-08-20.md); by the evening of the day it was written, three builder dispatches had each exceeded three times its stated ceiling. Meanwhile `analyst` ran **zero times** — its only run all-time resolved 2026-08-19 — on a day when the fund's binding constraint was a corrupted risk reference and its idea funnel produced no candidates. The bench is not underused; it is unevenly used, and the direction of the imbalance runs away from generation, which is leg 2 of the metric.

Three, small and worth fixing cheaply. The record keeps two clocks. Two artifacts produced by runs that resolved on 2026-08-20 UTC are filed under 2026-08-21 filenames — `docs/pm/PM_SLEEVE_V2_2026-08-21.md` (resolved 19:40:24Z) and `docs/research/QUANT_ENTRY11_2026-08-21.md` (resolved 20:13:36Z) — and the INTC termination written at 2026-08-20T21:39:27Z says in its own reason text "Closed by the CTO at the CEO's explicit direction, 2026-08-21" (seq 560). Separately, five agent runs share the identical `resolved_at` of 2026-08-19 22:54:53, which is a backfill, not five simultaneous resolves. None of this moved a dollar. All of it means that anyone reconstructing a day by filename or by resolve-timestamp will silently lose work that belongs to it — which is the exact failure this seat was created to end.

CTO FILING NOTE (2026-08-21, AT RESOLVE)

Spot-verified before filing, all exact: the 103 CEO decisions and the 34-item burst (19:54:56→20:01:52Z, 6m56s); the three guard-refusal citations (seq 373/374/395) word-for-word; the 3,216,188-token sum, which

independently matches the CTO's own audit sum computed hours earlier while judging the local-model audition; the 99-commit count, which reconciles with the day-pack. Donna also REFUSED two claims supplied in her own dispatch brief — "three usage-limit interruptions" (uncited anywhere) and the north-star metric as an 08-20 decision (the commit is dated 08-21) — which is the seat's hard rule ("the brief is a pointer, the log is the source") working against its own principal on day one. Her two record-keeping findings (UTC/local filename drift; the 2026-08-19 22:54:53 resolved_at backfill) and three org observations are recorded on run-secretary-1 for the CEO's read. This file is the first entry in docs/archives/; the PDF render follows via scripts/archive_pdf.py.