

...

TL;DR

This replaces the interim archive for 21 August, which was cut at midday by mistake and missed the entire second half of the day.

The book barely moved — NAV \$1,885.74, up 95 cents — but the policy that lets the machine sell without asking you was rewritten, because we found it would sell shares we do not own on 8 September.

Three pieces of work were killed by review and one survived a million adversarial tests; three candidates reached the belt and none passed.

Nineteen agent runs, 4.2 million tokens, and you accepted 78 items — seventeen of them in four and a half minutes.

Waiting on you: thirteen recommendations. Waiting on the chair: twenty-eight approved requests nobody has run yet.

...

THE DAILY · 2026-08-21

> **Supersedes the interim archive of the same date** ([eb9fa43](#)), cut at **12:45:20Z**

> because the chair read a local clock as end of day. It covered 99 of the day's 375

> events and said so. Nothing in it was false; it was half a day, and the completing

> half was owed. This is the full UTC day, re-derived from the record.

NAV	\$1,885.74 (+\$0.95 on the day)	Cash	\$968.69 (51.37% of NAV)
Gross	\$917.06 (48.63% of NAV, 4 legs)	Drawdown	1.17% of the 10.0% limit
Halt headroom	\$168.46	Dated items	5 — nearest 2026-09-05
Awaiting the CEO	13	Fills today	2 (\$414.77 notional)
Dispatches	19 resolved; ≥2 unrecorded	Defects confirmed	11

What moved

— **The auto-approval envelope went v3 → v4 and is live** ([b05cb9b](#) , 22:58:20Z): v3

passed all four live exits **twelve checks out of twelve** while the broker holds

zero of the underlying — **\$750.35** armed, 39.79% of NAV, **\$501.58** date-certain

on 2026-09-08 (DAY_LOG 2026-08-21).

— **The adversary returned its first SURVIVES** on that change — 1,067,152 cases, zero

orders v4 approves that v3 refused, 163× tighter — and **killed two builder diffs** the

same day (run-adversary-d8, run-adversary-d9).

- You accepted 78 recommendations and rejected none; the largest burst was **17 in 4m23s** at 18:54Z (`fund_events` , actor `ceo`).
 - Three candidates reached the belt at 20:08–20:14Z and all three failed — leg 2 off zero for the first time in two days (`fund_candidates`).
 - The book stopped matching the broker: 50 `ReconciliationMismatch` events from 15:21:39Z across ten symbols, broker holding 0 of every one; \$129.59, 6.87% (seq 715–854).
 - Two constitutional changes: decisions are now provisional and challenging one is a duty; the COO trigger moved 20 → ≥50 over the COO's recorded objection.
- Awaiting you (open state at 2026-08-22T00:00:00Z — 13 recommendations; 0 pending orders, 0 requests awaiting your approval)
1. Accept or reverse the desk-counter change (run-builder-d9 rec 1) — a loosening: the COO trigger fires later. The builder asked for explicit sign-off.
 2. Decide the CEO-page restructure (run-builder-d9 rec 2) — specified, not shipped.
 3. The gate's cost-robustness door (run-validator-breakeven recs 1–2, \$500 sleeve) — the criterion has never produced a number in 40 candidates.
 4. A challenge to your own excess-return amendment (run-validator-breakeven rec 5) — the seat flagged it as a loosening itself, so it goes to the adversary blind first.
- *Ninety-five cents, Neelesh — and a policy that would have shorted you in September, caught three weeks early by the people you pay to disagree with each other.*

THE RECORD • 2026-08-21

Full UTC day; cut 2026-08-22T00:00:00Z, written 2026-08-22. Sources: `fund_events` seq 563–937 (375 events), `fund_agent_runs` (19 resolved), `fund_candidates` / `fund_lean_sweeps` , `/fund/desk` , `/fund/risk/monitor` , `/fund/exits` , three git logs, `DAY_LOG.md` , `CTO_REVIEW_QUEUE.md` , 17 filed docs. Where a closing number differs from `eb9fa43` , the difference is the six hours it did not see.

NAV & THE BOOK

	OPEN — SEQ 574, 00:50:15Z	CLOSE — SEQ 844, 20:39:12Z
NAV	\$1,884.79	\$1,885.74
Cash	\$1,383.46 (73.40%)	\$968.69 (51.37%)
Positions	\$501.34 (26.60%)	\$917.06 (48.63%)
Legs	2 — DBC, TLT	4 — SPY, DBC, TLT, DBA
NAV per unit	0.942396	0.942871

Halted	false (resumed 00:03:25Z, seq 564)	false
--------	------------------------------------	-------

15 strikes; day change against 2026-08-20's last strike (seq 562) **+\$0.95**. The interim closed at \$1,884.79 / \$916.11 on seq 655 — right for its cut; the marks first moved at 15:21:39Z (seq 714).

Drawdown. Reference **rebased to \$1,908.09** at 13:07:57Z (seq 694, [neelesh-via-co-cto](#)), repairing a peak corrupted by the 2026-08-20 phantom GLD fill: 0.424471 GLD sold at a fabricated \$100.00 against a same-day close of \$415.04, **\$128.26** destroyed. [unrebased_peak_nav](#) preserved at \$2,036.35, [peak_basis: rebased](#). Closing drawdown **1.17%** of a 10.0% limit; halt line \$1,717.28, headroom **\$168.46** (NAV – 0.90 × rebased peak). No alarms open.

Leg 3. \$414.77 into risk, gross 26.60% → 48.63%. Cash sits \$874.45 above the 5% floor (\$94.29) — recorded as **not** a leg-3 defect because it carries a written reason: phase 1 at throttle, phase 2 dated (TRIAGE3).

|| TRADING & EXECUTION

SYMBOL	SIDE	QTY	PRICE	VENUE AS RECORDED	APPROVER	CITATION
SPY	buy	0.346119	\$762.5999755859375	paper	neelesh	seq 589 → 591
DBA	buy	5.314306	\$28.3799991607666	disputed, below	neelesh	seq 592 → 594
SPY	buy	0.346119	—	paper	worker — expired	seq 572 → 576
DBA	buy	5.314306	—	alpaca	worker — expired	seq 573 → 577

Both fills carry [approver: neelesh](#) — your click, not the policy, which fired zero times today. Both were staged 00:08Z, expired on the 120-minute staleness limit, re-staged 06:48Z, filled within three seconds. Exit rules were committed 00:06–00:07Z, ****before either order existed**** (seq 565–568).

The DBA venue, resolved against the interim: **"THE DBA LEG LABELLED 'alpaca' FILLED ON THE PAPER VENUE. Venue is not a route — one connector chosen at import"***

(run-riskofficer-2). The CEO-authorized R15 cost experiment produced ****zero information and had been marked done; the chair verified all three proofs and reopened R15****.

50 ReconciliationMismatch events, 15:21:39Z–20:39:13Z (seq 715–854): DBA, DBC, GLD, INTC, MSFT, NVDA, SOFI, SPY, TLT, XLE — five each, every one [actual: 0](#) against a non-zero book. Divergence \$129.59, 6.87% (seq 725).

||| RESEARCH & VERDICTS

Nineteen runs. The morning six are in [eb9fa43](#) §III and are not restated.

SEAT / RUN	VERDICT	ARTIFACT
riskofficer r-2	**SEVEN FINDINGS, two critical, and the largest is NOT in the brief** — the DBA venue; the rebase-direction fix needs two lines (the confirm echo hashes the same wrong value); risk-limits and trading-resume have no guard at all	AUDIT_RISKOFFICER_2
pm divergence	**FENCE THE COHORT... zero trading** — the fund has no instrument that writes to the venue without writing to the ledger. Unasked-for finding: **autopolicy.py has ZERO venue checks**	PM_DIVERGENCE
adversary D8	KILL , three grounds: **THE SIX FLAGGED LINES WERE CLEAN; the real widening was elsewhere** — <code>_real_broker()</code> → <code>_broker_is_real()</code> flips one of eight flag combinations refuse→allow, and that cell IS the live .env	ADVERSARY_BUILDER_D8
analyst cycle 3	5.02 drift DEAD (event t=-2.33, calendar-time t=0.65). Larger, chair-verified: **a ~44%/yr phantom low-minus-high price factor (t=4.62, positive all 7 years)**	ANALYST_CYCLE3
validator gate v5 r5	Financing fixed and unreopenable (zero-skill cash mix 0.0000%/yr at every weight, 0.0% in 16/16 MC cells, against +35%/yr in r4). **THE PREMIA RULE IS NOT ADOPTABLE** — discrimination 0.62, CI [0.53, 0.72] . CLOSED, not adopted	GATE_V5_ROUND5_MEASURED
coo triage #4	31 → 14 rows → 7 batches, ranked by reversibility not size . Challenge #1 **SUSTAINED IN DIRECTION, REFUTED IN EVIDENCE**; **the counter rebuilt 0 → 31 in one day from six seat runs**	TRIAGE4
riskofficer r-3	v4 specified, strictly tightening, three checks. **v4 MUST NOT SHIP ALONE — a v4 decline is currently invisible and a declined exit dies in 120 minutes and never re-raises**	R19_ENVELOPE_V4_SPEC
mechanism cycle 3	**0 candidates... the bottleneck has moved from idea supply to the instrument.** 4 of 8 verdicts across three cycles died on the device. Menu 15 → 19, zero containers	MECHANISM_CYCLE3
builder D9	Gates PASS (CH 1,324; KP 255/255), not merged — **the seat flagged its own diff as touching a control despite a green classifier, and the chair agrees**	two bundles
adversary D9	KILL , one narrow ground: **Server says 1, page says 0, ON THE SAME LINE OF THE SAME SCREEN.** All else held on 105 live rows; nothing auto-closes	verdict only
validator breakeven	**ZERO OF 40 — A CENSUS, NOT A SAMPLE.** Brief refuted by measurement: idle cash earns exactly 0.000% (n=11); what inflates is total return embedding each held asset's risk-free component — 10.4–18.4 bps against a 10.0 floor	verdict only

The v4 chain, and a hole in the run record. Specified by the riskofficer (21:14Z), built in an isolated worktree, reviewed **adversary blind** — **SURVIVES. 1,067,152 adversarial cases — 817,152 at `evaluate()` level and 250,000 end-to-end... ZERO orders that v4 approves and v3 would have refused** — merged `b05cb9b` 22:58:20Z, 1,323 tests green, fifteen injected mutations and fifteen caught. **Neither the builder run nor the adversary run has a row in `fund_agent_runs`**; the chair self-reported it — **shipping R19 without recording the run, so the CEO could not see it** (DAY_LOG). The day's largest control

change is evidenced only by a commit body.

Leg 2 — three candidates reached the belt, all failed: `a663a592ff1d` (20:08:26Z), `01593c65a05d` (20:11:17Z), `01b61967c933` (20:14:03Z), `passed = False` against a 41.55% benchmark; 15 sweeps. These are the analytics-capture re-runs (37|0 → 40|3). Under the Clean Field Rule **as amended the same day**, they are three NEW measurements on a NEW window and must never be set beside the fenced 37.

IV DECISIONS & GOVERNANCE

183 `DeskRecommendationDecided` events — CEO 78, co-CTO 69, CTO 36; accepted 86, done 63, staged 33, open 1. **All 78 CEO events were accepts; no rejections.** Bursts: 5 (166s), 5 (57s), 1, **18 in 709s** (12:41Z), **17 in 263s** (18:54Z), 11 (250s), 8, 8, 5 (28s).

DECISION	BY	WRITTEN REASON, AS RECORDED
Envelope v3 → v4 adopted, merged	CEO; chair on single-item Tier-3 authorisation	""get it done but make sure fable gets full context of the change""
Drawdown reference rebased to \$1,908.09 (PM R1)	CEO, staged <code>neelash-via-co-cto</code>	""Repair — rebase to \$1,908.09"" (seq 694)
Decisions are provisional — challenging one is a DUTY, admissibility bar, adversary pass for anything loosening	CEO	""my approved decisions needs to continually evolved and updated""
COO trigger >20 → ≥50	CEO	""Lets run coo on >=50 items or we can trigger as needed"" — 11 of 20 open recs already executed
— objection preserved	COO	""the number is not the defect, the blind spot is"" — refuted in evidence by the same seat eight hours later
Two agents in parallel when independent	CEO	five-part dependency test into the constitution
## <code>BINDS</code> protocol	CEO	""agree lest fix"" — a lesson that stays where it was found improves nothing
Harness phase gets an exit condition — five <code>alpaca-prod</code> preconditions	CEO	"fix the harness first" has no natural stopping point
Three fund modes; mock isolated, not ephemeral	CEO	""every order needs to route to alpaca paper account no sham and kill fake_firestore""
200-name universe FENCED ; Clean Field Rule amended to remove the re-run recovery path	CEO, on the quant's challenge	a re-run creates a NEW candidate on a NEW window
Closing a dispatch is the chair's judgement, never automatic	CEO	""no it should nto close automatically""

Also: gate v5 round 5 taken on CEO override of its Tier-3 parking; and ""four desk

requests approved as `neesh-via-co-cto` **, the first chair-approved requests in the fund's history, on `"this says awaiting you when its already accepted"` — awaiting `4 → 0`.

Three challenges on the rule's first day. #1 (chair's) refuted in evidence by the COO; #2 (COO's) accepted and the constitution amended the same day — **17 of 19 register entries carry a review trigger no code evaluates while the endpoint reports `triggers_unchecked: []` **; #3 and the mechanism's went to the CEO's desk.

V INSTRUMENTS & INFRASTRUCTURE

Commits: 112 UTC (ClarkHarness 49, KryptonPay 13, firm 50) against **101 local IST** (46 / 18 / 37); method `git log --all` on `%CI`, converted once each way. 2026-08-20 counted 99 local; the interim counted 83 local / 46 UTC at midday.

- **Envelope v4** `b05cb9b` — autopolicy tests 19 → 59 plus a wiring file for the tick, which had none; `context_for` had zero tests before this. Merge gate blocked on sensitive surface, which is the correct verdict and the route to the adversary.
- **Builder D8** — KryptonPay merged, ClarkHarness held then **killed**; round 2 repaired all three grounds and returned to the adversary. **D9** — both gates PASS, **killed**.
- **Offsite durability restored** (queue ~14:30Z, CEO-instructed): the fund had none.
- **17 docs filed**, all new. `docs/README.md`, named in the constitution as the carrier of each doc's status, was last committed **2026-08-20T14:26:17Z** and indexes none.

VI THE DEFECTS LEDGER

Leg 1 — eleven confirmed defects, four critical. The seven in the interim are not restated.

DEFECT	FOUND BY	MONEY IT COULD HAVE TOUCHED	CITATION
v3 approves exits against a broker holding zero — 12/12 checks pass, all four legs	riskofficer	\$750.35 armed, 39.79% of NAV; \$501.58 date-certain 2026-09-08, shorting enabled	r-3; <code>b05cb9b</code>
<code>autopolicy.py</code> has ZERO venue checks — one comment at line 14	pm, chair-verified	same exposure, second route	run-pm-divergence
A v4 decline is invisible — no event, no line; dies at 120 min, never re-raises, two docstrings claim otherwise	riskofficer	turns a silent short into a silently dropped exit	<code>exitrule.py:275</code>
P&L ignores the position's sign — on a short, a rising price reads as a gain	chair, from r-3	gates any short; open	<code>riskmonitor.py:878</code> ; <code>34338ef6</code>
The "alpaca" deployment filled on paper, marked done	riskofficer	R15's learning goal: zero cost information from \$150.82	run-riskofficer-2

DEFECT	FOUND BY	MONEY IT COULD HAVE TOUCHED	CITATION
USE_FAKE_FIRESTORE routes ORDERS, not just the ledger — three paths to PaperConnector , two silent	riskofficer + chair	the chair had flipped it	fund.py:151-163 ; b72847bc
store_backend() defaults to firestore when FUND_STORE is unset — the fund came up on the wrong store mid-restart, every read correct	chair, on itself	caught only by a 503 on a desk write	events.py:218
Cost-robustness has never produced a number — 40 candidates, 114 sweeps, 0 measured; all three belt passes ever sit in the 11 satisfied by absence	validator	the \$500 sleeve; bias 10.4–18.4 bps against a 10.0 floor	run-validator-breakeven
A ~44%/yr phantom price factor (t=4.62, positive all 7 years) — today-anchored splits plus a zero-attrition universe	analyst, chair-verified	every walk-forward verdict the gate has issued	run-analyst-cycle3
Belt capacity decided by an unseeded hash on a fill-count tie — 16.7× swing	quant, carried	a gate criterion	8c72939e
The sensitive-diff classifier is a regex and missed the only real widening	adversary	the live .env cell, on a ledger-writing endpoint	D8; d1d5beef

FOUR NUMBERS FOR ONE HAZARD, REPORTED AND NOT RESOLVED. Armed reads **\$750.63** (pm,

16:01Z), **\$750.36** (coo 20:19Z, riskofficer 21:14Z), **\$750.35** ([b05cb9b](#), 22:58Z).

Date-certain reads **\$502.15**, then **\$652.09**, then **\$501.58** — the \$652.09 corrected by the builder and chair-verified against [/fund/exits](#) as "summing two different dates" (DAY_LOG). Four re-derivations by four seats, converging but not converged.

A second disagreement. [CTO_REVIEW_QUEUE.md](#) states "61 [ReconciliationMismatch](#) events now exist at seq 749–807". The log holds **12** in that range, **50** on the day, **71** all-time (from seq 119, 2026-08-15). Reported, not resolved.

VII THE FLOOR

19 runs, 4,197,239 tokens, 1,484 tool uses, every one on Opus.

SEAT	RUNS	TOKENS	SHARE	SEAT	RUNS	TOKENS	SHARE
builder	3	1,900,254	45.27%	secretary	2	278,147	6.63%
adversary	3	439,525	10.47%	coo	2	264,881	6.31%
validator	2	335,367	7.99%	mechanism	2	259,689	6.19%
riskofficer	2	306,039	7.29%	pm	1	132,098	3.15%
analyst	2	281,239	6.70%	quant	0	—	—

Builder share **69.40%** → **45.27%** against the interim's midday read: nine non-builder runs landed after the cut. Mean cost per run 220,907 tokens (300,970 at the cut; 201,012 on

2026-08-20). **Graduated**: the desk API takes full ids and does not validate — an 8-character display prefix appends a resolution against a nonexistent aggregate, and six inert events stand as proof.

VIII CARRIED FORWARD

- **2026-09-08** — TLT and DBC time exits, \$501.58 date-certain. v4 refuses them, which converts a silent short into a **silently dropped exit** until the skip-visibility half ships. Undated `loss_pct` rules on all four symbols (TLT 4.0%, DBC 8.7%, SPY 7.3%, DBA 6.1%) make it a tomorrow risk, not a September one.
- **The sign-inverted P&L** (`34338ef6`) gates any short and is open.
- **Builder D10** — dispatched 2026-08-21, `**returned 2026-08-22T00:34:22Z`, merged `b5e15d5` / `b80a7290` at `02:25:35Z**`. Both belong to the 2026-08-22 record, as does its desk measurement (14.74 screenfuls → 0.47 before the first decision control). Noted here because the D9 kill it repairs is a 2026-08-21 fact.
- **Gate v5 round 5 CLOSED, not adopted**; round 6 inputs filed (`4698dee7` , `8e2c799a`). No risk-free series exists anywhere in the gate path.
- **Tier-3 parked for Fable, in the chair's order**: the register's trigger-evaluability fix (**before** registering governance decisions), guard v1.3's server-issued echo, risk-limits and trading-resume onto the approval channel, the rebase direction pair, D5, D7.

IX THE OBSERVER'S NOTE — AND THE FRICTION LEDGER

1. The desk did not fill because you were slow; it filled because the bench was fast.

78 accepts, no rejections, twice at one every fifteen seconds; thirteen open at midnight. The COO measured it from the other side: `"the counter rebuilt 0 → 31 in one day from six seat runs, so the trigger is calibrated against BENCH OUTPUT VOLUME, not CEO load."` The threshold it is judged against moved to ≥ 50 the same evening.

2. Every seat but the quant ran, and the builder's share halved — 45.27% against 69.40% at midday and 57.4% on 2026-08-20. **Leg 2 moved off zero**: three candidates, all failed. `docs/README.md` still indexes nothing — third day, unchanged.

3. WHO IS WAITING ON WHOM. At midnight `**28` desk requests stood `approved` and undispatched. `** Approved is not answered: a request is approved when a human agrees it should happen and answered when a seat runs.` **All 28 were waiting on the chair**, because a seat cannot start itself. Three have been answered since (`920ecbe5` , `907ecc74` , `af279b4c` , all 2026-08-22); the rest stand.

WHO ASKED	OF WHOM	WHAT	FILED (Z)	AGE AT MIDNIGHT	LAST MOVEMENT
mechanism	pm	close R8 formally in the sleeve record	09:26	14h 34m	approved 12:46, never dispatched
co-CTO	pm	sleeve phase 2 — dated 2026-09-05	12:54	11h 06m	approved 18:53
co-CTO	analyst	full-universe corpus deepening (your budget call)	12:54	11h 05m	approved 18:48
co-CTO	builder	remove <code>USE_FAKE_FIRESTORE (09e49ae5)</code>	14:28	9h 32m	approved 18:46
co-CTO	builder	spine wedged; liveness cannot report it	18:24	5h 35m	approved 18:46
co-CTO	builder	classifier missed the widening (<code>d1d5beef</code>)	18:32	5h 27m	approved 18:45
22 further rows	builder ×19, validator ×2, mechanism ×1	—	19:28–23:30	29m – 4h 32m	all approved within 30 min of filing

Two rows where the chair owed you, not a seat. (i) The v4 build and its adversary review were merged with **no run recorded** — *"shipping R19 without recording the run, so the CEO could not see it"* (DAY_LOG); the day's largest control change was invisible on your desk. (ii) You asked at 12:39Z for Donna's memo on the CEO desk (`920ecbe5`); the route was a **live 404** for the rest of the day, so the card rendered a permanent absence.

Two things nobody owns, and the fix for each is a person, not a build. The COO's ≥50 objection sits in the constitution marked preserved-unresolved — and the same seat **refuted it in evidence** eight hours later, so an objection and its refutation now sit side by side with nobody assigned to close either. And 17 of 19 register entries carry a review trigger no code evaluates while the endpoint reports `triggers_unchecked: []`.

THE LENGTH, DECLARED RATHER THAN OVERFLOWED SILENTLY. This record runs ~18,700 characters against a five-page limit of ~15,000 — roughly 6.6 pages. Nine items are already compressed to their citations to get here: the morning six runs (`eb9fa43` §III), offsite durability, guard v1.2, the four `neesh-via-co-cto` approvals, the six orphaned resolution events, the D8 round-2 return, the COO memo format, FLOW_TEST_SYNTHESIS, and the 17-doc list. The remaining 1.6 pages are the eleven-row defects ledger, the eleven-row verdicts table and the eleven-row governance ledger, and cutting any of them would cut a number rather than a sentence. The log has the detail.