

...

TL;DR

The book did not move in value and moved a lot in shape: NAV held at \$1,884.79 while \$414.77 of idle cash finally bought two new legs. Four seats said a well-argued NO — the gate's fourth kill, one dead study, one retired idea — and a builder refused a false instruction its own brief demanded, which is the save of the day.

Nothing reached the belt for the second day running; that is the number to watch.

Waiting on you: the drawdown-reference decision (~\$400 of blocked capacity), two COO batches, and four requests. Nothing is urgent within the hour.

...

THE DAILY · 2026-08-21

The book. NAV \$1,884.79 at the close (seq 655, 12:07:10Z), unchanged to the cent from the open (seq 574) — cash \$968.69 (51.40%), positions \$916.11 (48.61%) across four legs (SPY, DBC, TLT, DBA), halted **false** since 00:03:25Z (seq 564).

What moved

- Two fills, both approved directly by the CEO: SPY 0.346119 @ \$762.5999755859375 and DBA 5.314306 @ \$28.3799991607666, \$414.77 deployed, gross 26.60% → 48.61% (seq 591, 594).
- The adversary killed gate v5 round 4 — the fourth kill in the chain — on four grounds, and the CEO signed the fix: premia is judged on **excess** returns, financing charged (seq 632; commit b331114).
- The analyst declared entry 8 a **measured NO-GO** on 16,466 events and refreshed the corpus 863 → 1,035 observations; the mechanism declined entry 7, retired entry 4, and grew the menu 12 → 15 (run-analyst-cycle2, run-mechanism-cycle2).
- The builder **refuted a factual claim in its own dispatch brief** — the EDGAR "-4h" timezone shift is false — and the co-CTO confirmed it independently on 4,895 filings; applying it would have manufactured the lookahead the column exists to detect.
- Dispatches 6 and 7 both merged green (1277/1277 and 215/215), shipping the CEO's four desk queues, this archive's shelf, and the gate data path (commits fda19b9, c209bod).
- COO triage #3 took 23 items to 5 batches and recorded two dissents against the co-CTO chair — both accepted (run-coo-3).

Awaiting you (open state at 12:45:20Z; 27 items, 23 open recommendations + 4 requests)

1. **PM R1 — the drawdown reference.** \$128.26 already destroyed, ~\$400 of capacity blocked, headroom \$52.08 → \$167.51, halt odds 58.4% → 9.6% (COO Batch 1, accepted 12:43:22Z; the decision itself is still untaken).

2. **COO Batches 2 and 4, and rec 6** — still **open**. Batch 2 carries the 2026-09-08 auto-exit of \$501.34 that needs no click while its re-buy needs you.

3. **Four requests:** **dc7b068c** riskofficer audit, **27957634** pm closes R8, **909c316c** analyst 5.02 measurement, **920ecbe5** builder puts this memo on your desk.

4. **20 open recommendations** from analyst cycle 2 (5), mechanism cycle 2 (5), builder D6 (6) and D7 (4).

Two chairs, seven seats and 99 events to move NAV by exactly \$0.00 — though \$414.77 did finally go to work, Neelesh, which is more than the belt can say today.

THE RECORD · 2026-08-21

****Cut at 2026-08-21T12:45:20Z. Sources:** **fund_events** seq 563–661 (99 events), **fund_agent_runs** (7 runs resolved), the live spine, three git logs, and the filed docs. Every derived number below was computed from those sources and re-checked.**

NAV & THE BOOK

	OPEN — SEQ 574, 00:50:15Z	CLOSE — SEQ 655, 12:07:10Z
NAV	\$1,884.79	\$1,884.79
Cash	\$1,383.46 (73.40%)	\$968.69 (51.40%)
Positions	\$501.34 (26.60%)	\$916.11 (48.61%)
Legs	2 — DBC, TLT	4 — SPY, DBC, TLT, DBA
NAV per unit	0.942396	0.942396
Units outstanding	2,000.0	2,000.0
Halted	false (resumed 00:03:25Z, seq 564)	false

Ten NAV strikes were recorded (seq 574 through 655). The day change against the last strike of 2026-08-20 (\$1,884.79, seq 562) is **\$0.00**, and the reason is worth stating plainly: every mark in the close payload is bit-identical to the same mark yesterday (DBC \$31.110000610351562, TLT \$82.33999633789062), and the two new legs are held at the exact price they filled at. The \$0.00 is unchanged marks, not offsetting moves.

Close positions: SPY 0.346119 @ \$762.5999755859375 = \$263.95; DBC 8.122157 @ \$31.110000610351562 = \$252.68; TLT 3.019871 @ \$82.33999633789062 = \$248.66; DBA 5.314306 @ \$28.3799991607666 = \$150.82. Sum \$916.11.

Leg 3 of the team metric — capital deployed under mandate. \$414.77 moved from cash into risk today, the largest single-day deployment the book has recorded: gross rose 26.60% → 48.61% against a throttled target of 48.44%, 17bp over and stated in the order rationale rather than hidden (seq 572). Cash of \$968.69 sits \$874.45 above the 5% floor (\$94.24); the COO recorded this as **not** a leg-3 defect because it carries a written reason — phase 1 is at the throttled target and phase 2 is dated 2026-09-08 (docs/coo/TRIAGE3_2026-08-21.md; docs/pm/PM_SLEEVE_V2_2026-08-21.md). Drawdown at the cut: peak_nav \$2,036.35 on a trailing_365d basis, limit 10.0%, no alarms open (`GET /fund/risk/monitor`). Computed from those: drawdown 7.44%, halt line \$1,832.72, headroom **\$52.08** — reproducing the COO's independent 12:23Z read exactly.

|| TRADING & EXECUTION

SYMBOL	SIDE	QTY	PRICE	VENUE (AS RECORDED)	APPROVER	CITATION
SPY	buy	0.346119	\$762.5999755859375	paper	neelesh	seq 589 → 591
DBA	buy	5.314306	\$28.3799991607666	see note	neelesh	seq 592 → 594
SPY	buy	0.346119	—	paper	worker (declined)	seq 572 → 576
DBA	buy	5.314306	—	alpaca	worker (declined)	seq 573 → 577

The approver column is the governance story and it is clean: both fills carry `approver: neelesh` — the CEO's own click, not a channel identity and not the auto-approval policy. Guard v1.2 shipped the `neelesh-via-co-cto` identity today (commit d7a04bf) and **no event in the log yet carries it**, which is why the COO scoped that audit leg as a design review rather than an evidence review (run-coo-3 rec 1). Both orders were staged twice. The 00:08Z pair expired on the 120-minute staleness limit before the CEO woke (seq 576, 577, approver `worker` : "the price and signal behind it have moved"); they were re-staged at 06:48Z with the CEO present and filled within three seconds of approval. Exit rules were committed at 00:06Z and 00:07Z — ****before either order existed**** (seq 565–568): SPY loss 7.3% / time 2026-11-19, DBA loss 6.1% / time 2026-11-19. Three sleeve strategies were registered with `claim_type: premia` at 00:08Z (seq 569–571), including `sleeve_beta_500` retroactively after being absent from the register for three PM reviews.

A disagreement in the record, reported and not resolved. The DBA order was proposed to the **alpaca** venue as the fund's first experimental cost-measurement deployment (seq 588, PM R15). Its three lifecycle events do not agree: `OrderProposed` says `venue: alpaca` (seq 588), `OrderSubmitted` says `venue: paper` (seq 593), `OrderFilled` says `venue: alpaca` (seq 594). The fill's `avg_price`

(28.3799991607666) equals the recorded `arrival_price` to the final digit — which the shared API card documents as the paper venue's signature ("paper.py:116 fills at its own quote — the same call pipeline.py:215 records as arrival_price, so execution slippage is identically zero at any sample size"). Separately, neither of today's fills appears in `GET /fund/tca` at the cut: that endpoint's order list carries 20 orders and today's two (`867cabff`, `17d64dcd`) are not among them, both new sleeve ids are absent from `by_strategy`, and the informative sample still reads 8 execution splits against a stated reliability bar of 20. On the record as it stands, the learning goal R15 was deployed for is not yet advanced by a measured number.

||| RESEARCH & VERDICTS

adversary (run-adversary-v5r4, docs/reviews/ADVERSARY_GATE_V5_R4_2026-08-21.md) — **KILL**, the fourth in the gate v5 chain. Verdict quoted: "(1) Vol-matching levers TOTAL returns with no financing cost — the deterministic gift $(k-1)*rf$ beats the 2%/yr margin for any candidate at $\leq 60\%$ of benchmark vol; a zero-skill 40/60 SPY/BIL sleeve clears both full-sample legs on the fund's own bars. (2) $VR \leq 2.0$ guards dominance not presence: the same $\rho = .98$ wander diluted to 10% of idio variance passes 23.6%. (3) Headline FPR/TP measured in an unreachable 27-fold geometry; the shipped window_for caps hold-21 at 5 folds — reachable break-even prior 35.8%, worse than the 15.8% that killed r3. (4) No data path... SURVIVES: the four-leg structure (proven the real fix at $\rho=0$), the v5_alpha measurability fix, the taxonomy."*

analyst (run-analyst-cycle2, docs/research/ANALYST_CYCLE2_2026-08-21.md) — corpus **conditionally PIT-safe**, entry 8 **NO THESIS**. Quoted: "enter at the open of the session AFTER `filed`, never the close — 55.9% of filings accepted post-close on their own filing date; SRPT's 10-Q 106 seconds after the bell... Refreshed 51 tickers, corpus 863 -> 1,035 obs. ENTRY 8: NO THESIS — 16,466 events, 5.6y; pooled $t=2.22$ collapses to 0.85 ex-2021 and -0.27 since 2024, beta 1.27, ~\$38/position at sleeve size; independently reproduces the published replication collapse... The naive study would have shown $t=2.92$ and consumed a full candidate cycle — leg 1 protecting leg 2."*

mechanism (run-mechanism-cycle2, docs/research/MECHANISM_CYCLE2_2026-08-21.md) — entry 7 **DECLINED-WITH-CONDITIONS**, entry 4 **RETIRED**. Quoted: "ENTRY 4 RETIRED (11 of 12 configs lose; no convexity; ex-ante wrapper set SR +0.05/+0.21 with WTMF -7.8% in 2022; unjudgeable anyway — zero bear markets in any fold). DEFECT D4: Sharpe is leverage-invariant on excess returns, so long-only de-risking can never be premia — kills the whole family in advance. Menu 12 -> 15... Zero containers spent."* Both decline conditions are single-command measurables (C7a: $\text{corr}(\text{beta}, \text{excess Sharpe})$ turns negative,

today +0.62; C7b: anti-beta wrapper 3y excess turns positive, today -7.89%/yr).

coo (run-coo-3, docs/coo/TRIAGE3_2026-08-21.md) — **23 items → 12 needing judgement + 4 invisible decisions → 5 batches.** Quoted: **"Counter accurate this run (open-only fix holding), but 11 of 20 open recs were already executed. SHARPEST FINDING IS STRUCTURAL: PM R1 — the largest-money decision in the firm (\$128.26 destroyed, ~\$400 capacity blocked, halt odds 58.4% vs 9.6%) — carries status 'accepted' while the decision it asks for was never taken, so the counter cannot see it... SECOND: the 2026-09-08 TLT+DBC time exits (\$501.34) pass every auto-approval envelope check and will execute with NO click, while the re-establishment needs the CEO and nothing schedules him."**

builder — two dispatches, both merged (see §V).

secretary (run-secretary-1, docs/archives/2026-08-20.md + .pdf) — the first Daily, filed and rendered this morning. Five recommendations, all marked NOTED at resolve (seq 636–640); the notes-vs-suggestions contract was written into the seat definition the same day (commit 962144a).

Leg 2 of the team metric — candidates reaching the belt: ZERO. **fund_lean_sweeps** records 0 submissions today against 15 on 2026-08-20; **fund_candidates** records 0 started today against 3 on 2026-08-20 and 37 in total. Two funnel cycles ran and produced three well-argued negatives, one menu grown 12 → 15, and no candidate. This is the second consecutive day at zero.

IV DECISIONS & GOVERNANCE

DECISION	DECIDED BY	WRITTEN REASON, AS RECORDED
Premia judged on EXCESS returns, financing charged on leverage	CEO	**Under rf=0 the gate certifies T-bills as premia — that must be signed, not inherited from an omitted term** (seq 632; commit b331114)
COO triage trigger >20 → ≥50	CEO	verbatim: "Lets run coo on >=50 items or we can trigger as needed." Basis recorded: 11 of 20 open recs were already executed — "the counter was summoning the seat on stale bookkeeping" (commit 2d6813d)
— objection preserved beside it	COO (Vishesh)	"keep it at 20... The number is not the problem... The real fix is that the counter cannot see items like R1" — interest disclosed by the seat itself (TRIAGE3, review section)
Co-CTO chair created, three tiers, fail toward the queue	CEO	Fable's tokens run out; the firm must keep working in the gap without anything Fable built being reversed (constitution; commit ee94a8a)
Approval guard v1.2 ships neeleash-via-co-cto	CEO	distinct identity so the record shows which chair staged what (commit d7a04bf)
Third write exception: the secretary files her own archive + PDF	CEO	"she should generate the pdf not you from next run" (commit aa5b756)

DECISION	DECIDED BY	WRITTEN REASON, AS RECORDED
Secretary recommendation contract: note vs suggestion	CEO	""this seems more like a note and I don't know what to accept"" (commit 962144a)
Quant local split versioned in — qwen3.8 drafts sub-functions, Opus assembles	CEO	""agree on the quant layout"", on measurement: 4/4 hidden-test pass at ~102 tok/s (commit d61e504)
CIO trigger registered — audition when funded-candidate capacity binds	CEO	the CEO is the CIO today; at two sleeves the direction-setting load is small (commit 249bc2c)
No pushes until the CEO private workspace exists	CEO	standing (commit 4da6dfb)
Batch acceptance cascades ; validate-once before any sweep; v3 envelope text corrected	CEO	doc drift caught by the COO, ""a doc drift, not a loosening, since v3 is strictly tighter"" (commit 8ba83c3)
Artifacts file under the UTC date of the work they record	CTO	adopted in response to the secretary's day-one finding (seq 636)
Builder D7 merged, spine restarted, PIT backfill applied	co-CTO, Tier 2	Fable's handoff, verbatim: ""Merging on a PASS with 0 sensitive/forbidden surfaces is Tier 2"" (CTO_REVIEW_QUEUE.md)
API card's EDGAR line deferred , then quarantined	co-CTO, Tier 3 → COO dissent	""the API card is the CTO chair's instrument""; then, on the dissent: ""quarantining a false instruction is not the same as overruling a chair's judgement, and it takes seconds""
PM R1 un-gated from the riskofficer audit	co-CTO, on COO dissent	""accepted as arithmetic: I am correcting the handoff line to gate only rebase #2""

Decision volume. 51 `DeskRecommendationDecided` events: **15 accepted** (the CEO's clicks), 32 marked done, 4 staged. The CEO's accepts fell in four bursts — 5 in 166s (06:40–06:43Z), 5 in 57s (07:49–07:50Z), 1 at 08:13Z, and 4 in 125s (12:41–12:43Z, the COO batches). Three desk requests were approved, one declined for misrouting (seq 644: ""kind=audit forced serves=validator, but this is auto-approval-adjacent risk-engine supervision — the riskofficer lane""), and three resolved.

A disagreement between the constitution and the running instrument. The constitution now registers the COO trigger at ≥ 50 (commit 2d6813d). At the cut, `GET /fund/desk` serves `"threshold": 20`, `"coo_triage_due": true` against a `desk_load` of 27. The coded threshold has not moved. Reported, not resolved.

V INSTRUMENTS & INFRASTRUCTURE

Builder dispatch 6 (run-builder-dispatch6, 623,003 tokens, 398 tool uses) — merged as

`fda19b9` (ClarkHarness, +3038/-38) and `cbc32b8a` (KryptonPay, +2908/-35). Verdict:

""ALL BRIEF ITEMS + OPTIONAL ADDENDUM DELIVERED... ClarkHarness 1208/1208, KryptonPay 191/191 (corrected glob), tsc clean, forbidden surfaces o... Four live defects found by looking (equity-chart axis lie, archived-as-live, bench 0.0% fabrication, CEO desk blind

to its own asks) — all fixed."* It also shipped `scripts/merge_builder.py` (commit 63fe68a), the gate a bundle now passes before a human merges.

Builder dispatch 7 (run-builder-dispatch7, 838,995 tokens, 73 tool uses) — merged as `c209b0d` (ClarkHarness, +2309/-24) and `63454533` (KryptonPay, +1052/-187). Gates run by the co-CTO before merge: **ClarkHarness 1277/1277, 19 ordinary / 0 sensitive / 0 forbidden; KryptonPay 215/215, 14 ordinary / 0 sensitive / 0 forbidden.** Shipped: the CEO desk as four queues by person (Vishesh / Donna / Fable / Others), the secretary's floor desk and archive shelf, interaction durability (`fund_agent_transcripts` + a runs mirror), the EDGAR point-in-time column, the 8-K EX-99.1 exhibit follow, the crypto-namespace fix, floor run counts, and the gate data path (commit 76784c2, answering adversary r4 rec 4). Part G was deliberately not started at the sanctioned boundary. The documented completion step ran clean: **249/249 accessions resolved, 1,035 rows updated, 0 unresolvable**. The archive shelf is live: `GET /fund/desk/archives` returns 1 daily with its PDF at the cut. `fund_agent_transcripts` exists and carries **0 rows** — the table shipped today and has not yet been written to.

Commits. 83 by the local (IST, UTC+05:30) clock — ClarkHarness 37, KryptonPay 16, the firm repo 30 — against **46 by the UTC clock** (24 / 7 / 15). Method: ``git log --all --date=format-local:'%Y-%m-%d' , filtered on the date, run once under TZ=UTC`` and once under the local zone. The 2026-08-20 archive counted 99 by the local clock, so the local figure is the comparable one; the UTC figure is the one the firm's own filing rule now implies. Heads at the cut: ClarkHarness `30ed60c` , KryptonPay `a3d59de1` , firm `2d6813d` .

Docs filed today (10, all newly added): `docs/archives/2026-08-20.md` + `.pdf` ,
`docs/coo/TRIAGE2_2026-08-20.md` , `docs/coo/TRIAGE3_2026-08-21.md` ,
`docs/ops/FLOW_TEST_SYNTHESIS_2026-08-21.md` , `docs/pm/PM_SLEEVE_V2_2026-08-21.md` ,
`docs/research/ANALYST_CYCLE2_2026-08-21.md` ,
`docs/research/MECHANISM_CYCLE2_2026-08-21.md` ,
`docs/research/QUANT_ENTRY11_2026-08-21.md` ,
`docs/reviews/ADVERSARY_GATE_V5_R4_2026-08-21.md` . `docs/README.md` — which the constitution names as the carrier of each doc's status — was last committed 2026-08-20 19:56 local (`17d0cfd`) and indexes **none** of the ten, nor the archives shelf.

VI THE DEFECTS LEDGER

Leg 1 of the team metric. Seven confirmed defects in the fund's own beliefs, plus six found in the record while writing this.

DEFECT CONFIRMED	FOUND BY	MONEY IT COULD HAVE TOUCHED	CITATION
Vol-matching levers total returns with no financing cost — a zero-skill 40/60 SPY/BIL sleeve clears both full-sample legs on the fund's own bars	adversary	the gate is the only route to the \$500 alpha sleeve; adopting r4 would have certified zero skill as premia	run-adversary-v5r4
VR≤2.0 guards dominance, not presence — $p=.98$ wander at 10% of idio variance passes 23.6%	adversary	same instrument	docs/reviews/ADVERSARY_GATE_V5_R4_2026-08-21.md
r4's headline FPR/TP measured in an unreachable 27-fold geometry; reachable break-even prior 35.8% vs the 15.8% that killed r3	adversary	would have shipped a gate calibrated on folds the belt cannot produce	same
EDGAR <code>acceptanceDateTime</code> "-4h" is FALSE — the raw stamp is genuine UTC; the brief demanded the shift	builder (D7), confirmed independently by the co-CTO	all 1,035 corpus observations; the shift would have manufactured the sub-daily lookahead the column exists to detect	run-builder-dispatch7; CTO_REVIEW_QUEUE.md; builder n=2,400 + n=30,732, co-CTO n=4,895 across 6 issuers
A second drawdown rebase raises the effective peak (direction checked against <code>unrebased_peak_nav</code> , which never moves)	builder (D6)	\$0 today — never triggered; gates PM R1, which carries \$128.26 destroyed and ~\$400 of blocked capacity	<code>app/api/v1/fund.py:3619</code> ; request <code>dc7b068c</code>
D4 — Sharpe is leverage-invariant on excess returns, so long-only de-risking can never be a premia claim	mechanism	kills the whole family in advance; retired entry 4 and conditioned entry 7 at zero container cost	run-mechanism-cycle2
Items at status <code>accepted</code> whose execution needs the CEO personally are invisible to the desk counter	coo	three live, including PM R1 — the largest-money decision in the firm	run-coo-3; docs/coo/TRIAGE3_2026-08-21.md

Two harness defects were confirmed by the analyst and shipped the same day: the 8-K reader read cover pages only (83% zero-yield, item 2.02 earnings content unreachable) and `fetch_daily_bars('BTC')` returned CoinGecko bitcoin for an SEC filer (6 corpus observations) — both fixed in `2abadd2`. Entry 8 is the day's clearest honest negative: 16,466 events measured, pooled $t=2.22$ collapsing to -0.27 since 2024, and the naive version of the same study would have read $t=2.92$ and consumed a full candidate cycle. Found in the record while writing this Daily, each reported and none resolved:

1. `GET /fund/desk` serves `threshold: 20` while the constitution registers ≥ 50 (§IV).
2. The DBA order's three lifecycle events disagree on its venue, and its fill price equals its recorded arrival price exactly (§II).
3. Today's two fills are absent from `GET /fund/tca`'s 20-order list; the informative sample still reads 8 against a stated bar of 20 (§II).
4. `CTO_REVIEW_QUEUE.md` — the one artifact whose purpose is letting Fable reconstruct the gap — stamps all three of its entries `2026-08-21 ~18:10Z`. The commits they describe are `c209b0d` at **12:14:16Z**, `96ede33` at **12:21:28Z** and `2d6813d` at **12:40:22Z**; 18:10 is the IST wall clock. The label is off by +5:30 and collapses 26 minutes into one stamp.
5. `docs/README.md` indexes none of today's ten filed docs (§V).
6. `GET /fund/desk` seat telemetry reports `mechanism` and `builder` as `running_now: true` at the cut, while the same payload gives each a `last_run_at` hours earlier and both runs are resolved in `fund_agent_runs`. The dispatch events (seq 648, 649) have no matching completion event.

VII THE FLOOR

Seven runs resolved, **2,106,788 tokens**, 680 tool uses, every one on Opus.

SEAT	RUNS	TOKENS	TOOL USES	ARTIFACT
builder	2	1,461,998	471	two merges (D6, D7)
adversary	1	146,973	47	ADVERSARY_GATE_V5_R4_2026-08-21.md
analyst	1	136,223	47	ANALYST_CYCLE2_2026-08-21.md
coo	1	122,298	42	TRIAGE3_2026-08-21.md
mechanism	1	120,355	38	MECHANISM_CYCLE2_2026-08-21.md
secretary	1	118,941	35	archives/2026-08-20.md + .pdf
pm, quant, validator, riskofficer	0	—	—	—

Builder took **69.40%** of the day's tokens on 2 of 7 runs (1,461,998 / 2,106,788, cross-checked against `seat_telemetry.builder.tokens_today`). Mean cost per run rose to 300,970 tokens from 201,012 yesterday, on fewer than half the runs (7 against 16).

New seats and identities. The co-CTO chair was created today with a three-tier charter (constitution; `.claude/state/co-cto.md`; ledger `.claude/state/CTO_REVIEW_QUEUE.md`), and the CIO trigger was registered without seating a CIO — "audition a CIO seat the way the CDO was auditioned... never seat it by org-chart symmetry."

Workflow lessons that graduated. (a) Every gate table must run through the shipped `window_for` geometry or be labelled as measuring a proposed generator (seq 630, into `adversary.md` and `cto.md`). (b) The standing gate-audit battery gains a cash-mix null

- and a masked variant of every adversarial null (seq 631). (c) Artifacts file under the UTC date of the work they record (seq 636) — adopted from this seat's day-one finding.
- (d) `scripts/merge_builder.py` before every builder merge — used for both merges today.
- (e) The COO's accepted-item second pass is proposed as standing method (run-coo-3 rec 6, still open).

VIII CARRIED FORWARD

- **PM R1, the drawdown reference** — accepted as a batch, undecided as a question. Headroom \$52.08 → \$167.51 and halt odds 58.4% → 9.6% turn on it (run-coo-3 rec 1).
- **COO Batches 2 and 4, and rec 6** — three of seven run-coo-3 recommendations remain open. Batch 4's unlocking decision is stated as `"WHO DRAFTS round 5."`
- **Gate v5 round 5** is parked as Tier 3 for Fable. All four named prerequisites now exist: excess returns (signed), the shipped `window_for` geometry, the cash-mix and masked nulls, and the data path (commit 76784c2).
- **Four requests awaiting a human key:** `dc7b068c` (riskofficer, four-item batch), `27957634` (pm closes R8), `909c316c` (analyst 5.02 measurement — recorded as a measurement, never later citable as a proposal), `920ecbe5` (builder, this memo onto the CEO desk). `23b075a6` is approved and discharged by D7.
- **2026-09-08** — the TLT and DBC time exits close \$501.34 inside the auto-approval envelope with no click; the re-establishment needs the CEO and nothing schedules him. The COO asks for a PM dispatch on 2026-09-05.
- **The EDGAR correction is half-closed.** The card's false line is quarantined with a REFUTED banner; the replacement text is parked for Fable. The COO's condition stands: analyst rec 5 and builder D7 rec 1 close as a pair or not at all.
- **Unresolved for the riskofficer's first read:** the audit request `dc7b068c` cites `fund.py:3511`; on the current head the line reads at **3619**. Two merges moved it.
- **Tomorrow's first session** looks at PM R1 and the two open COO batches, in that order — they are the only items on the desk with money attached.

IX THE OBSERVER'S NOTE

1. The trigger worked today, and was retuned the same hour so that it would not have.

The COO trigger fired at `desk_load` 23 and the triage was dispatched at 12:22:34Z (seq 656), 16 minutes before it delivered. Compare yesterday, when this seat recorded the desk filling to 73 and the CEO clearing 34 items in 6m56s: today's largest single burst was **5 accepts in 166s**, and 15 accepts across the whole day. Standing observation (a)

did not repeat, and the reason is that the instrument arrived at a desk one-third the size. Then the threshold moved to ≥ 50 — a desk of 23 would not have summoned it. The written basis is sound and specific (11 of 20 open recs were already executed; the counter was firing on bookkeeping), and the CEO's stated fallback is manual dispatch at any count. But the COO disclosed its own interest and argued to keep 20 on a different ground entirely — that the counter is blind to **accepted** -but-undischarged items, three of them live today including the largest-money decision in the firm. The amendment does not touch that blind spot, and says so. Both facts belong next to each other in the record: the threshold moved in the loosening direction on a real measurement, and the defect the seat raised is still open. Meanwhile the running spine still serves 20.

2. Two seats spent 69.4% and 5.7% of the bench, and the belt stayed empty. The builder took 1,461,998 tokens on two runs; the mechanism took 120,355 on one. Standing observation (b) did not just repeat, it grew — 57.4% yesterday, 69.40% today. Standing observation (c) improved and is not yet fixed: the analyst ran once (zero yesterday) and the mechanism once, which is the funnel working exactly as designed — and the output was three honest NOs, a menu at 15 entries, and ****zero candidates on the belt for the second consecutive day**** (0 sweeps against 15 on 2026-08-20). Every one of those NOs was cheap and correct; D4 alone kills a family of ideas in advance and cost zero containers. The COO put it precisely, and it is worth repeating rather than softening: **"Every individual kill has been correct. Four consecutive kills while the funnel fills behind the gate is no longer a quality signal; it is the firm's binding constraint."** Four bench seats — pm, quant, validator, riskofficer — did not run at all. Leg 3 moved \$414.77 today, which is real. Leg 2 has not moved in two days.

3. The same defect class landed twice in one day, in the two places that record time.

The builder refuted a timezone instruction that would have shifted 1,035 stored stamps by minus four hours; the co-CTO confirmed it on 4,895 filings and called it the best call of the session. Hours later, the ledger written for the returning chair stamped its own entries **~18:10Z** when the commits they describe landed at 12:14Z, 12:21Z and 12:40Z — the local clock wearing a Z. And this seat's day-one finding on misdated filenames produced a standing UTC rule this morning (seq 636), while today's commit count differs by 37 between the two clocks (83 local, 46 UTC). None of these cost money. They cost the one thing this archive exists to protect: a reader's ability to reconstruct the order in which things happened. The root cause the co-CTO named for EDGAR applies to all three — **"a true measurement became a false instruction without anyone lying"** — and the fix it proposed, that every claim carries its verifying command, is the same fix a timestamp needs: carry the zone you actually read.

