

FIRST SUBSTANTIVE GATE PASS IN FUND HISTORY

The Announcement Premium

A scheduled-event liquidity strategy — the first real candidate to walk the entire research chain, mechanism to gate, and clear every criterion.


Gate v4.1 · PASSED · failures: none. Candidate 144387901688, judged 2026-08-22. Every prior pass in fund history was a deliberately planted null strategy used to expose the broken v1 gate. This is the first pass by a strategy that was meant to work — and the same run surfaced three defects in the gate itself, all candidate-flattering, all now filed.

+118.6%	+84.8%	80.4	1.33	\$19.9M
RETURN · 2.48Y · SIM	BENCHMARK · SAME BOOK	PSR · FLOOR 65	HOLDOUT RETENTION	CAPACITY EST.


The idea

Earnings arrive on a schedule. Each name's next 8-K (Item 2.02) date is projected +364 days from its last, and the book tilts into the **40 names nearest their announcement window** — collecting the premium paid by holders who de-risk ahead of scheduled news. The counterparty is the investor paying for certainty; the strategy is paid for holding through the event. 170-name universe, equal-weight residual, exit three days after the window. 21-day hold cadence, ~56× annual turnover.

Equity vs. its own benchmark



Growth of 100, 2024-02-26 → 2026-08-21. The shaded strip is a measured artifact: the benchmark's data ended 08-04 while the strategy ran to 08-21, over-crediting the reported excess by 11.85pp (35% of it). Like-for-like: **+21.98pp**.




Drawdown of the simulated book. Max -15.3%; the book runs at 0.66× its benchmark's volatility (beta 0.54).

The alpha claim, isolated




Cumulative return spread, strategy minus benchmark, both to the benchmark's own end date (like-for-like). Ends at **+22.0pp**. The honest statistic on session returns: +4.86%/yr on 12.66% tracking error — IR 0.38, $t = 0.60$. The gate's PSR was computed on the total book (Sharpe 2.31, $t = 3.60$), which is mostly the tape.

Out-of-sample, four folds



Walk-forward: train each fold, test out-of-sample, measure retention. 16 months of OOS union; fold 3 (Oct–Feb) is the regime that did not hold.

Cost fragility, measured by hand



The gate skipped its own cost check (a branch no candidate had ever reached). One extra 20bp container bracketed it: **breakeven 13.9 bps/side** — 1.4× the floor, not the 6.5× the gate implied.

READ IT THE WAY WE DO — THE PASS IS REAL, AND THINNER THAN THE HEADLINE

1. **The edge is not yet statistically distinguishable from zero** (active $t = 0.60$). The passing statistic rode a beta-0.54 book.
2. **Like-for-like excess is +22pp, not +34** — the benchmark-window artifact above, found and measured this run.
3. **All headlines struck at 1bp slippage** against the fund's 5bp default; at 5bp the excess is $\approx +15$ pp.
4. **Vol ratio measured 0.66** vs the 1.00 pre-committed — this behaves like a risk-premia harvest, and it passed the harder alpha gate anyway.

How it was earned — the chain, end to end

No stage was skipped, no agent occupied two stages of the same candidate, and the reviewer never saw the author's reasoning. Every number below is from the run record.

STAGE	WHAT ACTUALLY HAPPENED
1 · Proposal	The mechanism seat proposed the edge with a named counterparty and a pre-committed claim-type rule (vol ratio vs benchmark), filed falsifiable — what would prove it wrong, written before any backtest existed.
2 · Blind attack	The adversary seat received the proposal with the author's reasoning withheld. Verdict: SURVIVES as a belt candidate; its premia label KILLED by the candidate's own pre-commitment. It also issued a pre-run breakeven estimate — ~12 bps — that the belt later measured at 13.9. The estimate is now calibrated.
3 · Implementation	The quant seat translated the proposal into a LEAN algorithm — 170 names, scheduled-date projection, k=40 tilt engine — inside a sandboxed read-only container mount. The final run wrote zero new lines : the implementation stood as reviewed.
4 · The belt	22 LEAN containers, 96.4 minutes wall-clock : 5 cost sweeps × (3-point slippage grid + holdout leg), a full verification run, and one quant-initiated probe at 20bp to measure what the gate skipped. 6,997 fills; 8,662 portfolio targets; ~460–515s per container (170 sequential data legs each). Zero timeouts, zero failed jobs. Mid-run, the host session was restarted — the belt's own thread survived and completed unattended.
5 · The gate	Walk-forward geometry computed, not assumed (4 folds × 84-day OOS legs). Verdict: passed, failures none — and the run surfaced three candidate-flattering defects in the gate/belt itself (the skipped cost floor, the benchmark-window truncation, a capacity tie-break), each verified line-exact in code and filed for the next gate version.
6 · Disposition	Not deployed. No money moved. Parked for re-judgment under gate v5: excess returns, honest 5bp costs, like-for-like windows, active-return breakeven. A committed loss-stop (strategy-owned, predating entry) is required before any real stage.

mechanism → adversary (blind) → implementation → belt · 22 containers → gate v4.1 → **PASS** · For context: 40 prior gate verdicts in fund history, zero substantive passes before today.

What today proves is not this strategy — it may well die under honest costs, and the active t-stat says it should be presumed mortal. What today proves is the **process**: an idea walked every stage without a human hand on the scale, the machine told the truth in both directions, and the pass itself made the gate harder to fool. The next twenty candidates ride the same rails.